Prompt Collection for Is This a Time for Irony? Detecting Ironic-Response Affordance in Dialogues

Michael Oliverio, Pier Felice Balestrucci, Marta Marchiori Manerba, Luca Anselma, Valerio Basile, Cristina Bosco, Sergio Lopez Sancio, Viviana Patti, Erhan Sezerer, Antonio Uva, Alessandro Mazzei, and Davide Bernardi

Abstract

This document contains the complete prompts used for data augmentation, LLM-as-a-judge evaluation, ironic-response affordance detection, and perspectivist affordance detection. It is intended for inclusion in the project repository.

I. Prompts

This section reports all the prompts used throughout the CANTRIP corpus creation and experiments.

A. Data Augmentation Prompts

CSC Data Augmentation Prompt

```
Given the context and the reply of user U2 to Steve (whose relation to U2 is: {relation}):
- context: {context}
- U2: {reply}

Your task is to generate EXACTLY {num_turns} preceding dialogue turns (each turn
consisting of a message from Steve followed by a message from U2), PLUS the Steve message
that directly precedes the final provided U2 reply.

STRICT REQUIREMENTS:
- You MUST output exactly {num_turns + 1} turns total:
  - {num_turns} generated turns
  - + 1 final turn consisting of:
    - Steve: the message that directly precedes the final given U2 line
    - U2: the provided reply ({reply})
- All generated messages MUST be coherent, stylistically consistent, and logically lead to
the provided reply.
- Absolutely DO NOT invent or repeat any part of the provided reply.
- DO NOT output placeholders or generic fillers (e.g., "First message from U1").
- DO NOT repeat messages across different turns.
- DO NOT produce duplicated messages or self-repetitions across any turn.
- Preserve the writing style, tone, sentiment, and personality of U2.
- Messages from U1 must sound like Steve and must fit the relation: {relation}.
- The context must be respected and smoothly woven into the preceding turns.
- Output MUST be valid JSON ONLY. No comments, no extra explanations.

OUTPUT FORMAT (MANDATORY):
Return ONLY a JSON object in this exact structure:
<JSON>{json_structure}</JSON>

Produce only valid JSON. No explanations, no comments, no extra text.
```

MP Data Augmentation Prompt

```
You are given the following pair of final messages exchanged on social media between user
U1 and user U2:
- U1: post
- U2: reply
```

M. Oliverio, P. F. Balestrucci, M. Marchiori Manerba, L. Anselma, V. Basile, C. Bosco, V. Patti, and A. Mazzei are with the Department of Computer Science, University of Turin, Turin, Italy.

S. Lopez Sancio, E. Sezerer, A. Uva, and D. Bernardi are with Alexa AI, Amazon, Amazon Development Centre Italy, Turin, Italy.

```
Your task is to generate EXACTLY {num_turns} dialogue turns (each turn = one message from
U1 and one message from U2) that logically and stylistically lead up to the given final
messages.

STRICT REQUIREMENTS:
- You MUST output exactly {num_turns + 1} turns total: the {num_turns} generated turns +
the final provided turn.
- All generated messages MUST be coherent, consistent, and plausible precursors to the
final turn.
- Absolutely DO NOT invent or repeat text from the final messages.
- DO NOT repeat messages across different turns.
- DO NOT output placeholder or generic text such as "First message from U1".
- DO NOT produce duplicated messages or self-repetitions across any turn.
- Preserve the writing style, tone, sentiment, and personality of U1 and U2.
- Make the conversation flow naturally toward the final turn.

OUTPUT FORMAT (MANDATORY):
Return ONLY a JSON object in this exact structure:
<JSON>{json_structure}</JSON> Produce only valid JSON. No explanations, no comments, no
extra text.
```

B. *LLM-as-a-Judge Prompts*

CSC LLM-as-a-Judge Prompt

```
<TASK_OVERVIEW>
You are an AI judge model. Your task is to evaluate a generated dialogue based on a
specific dimension. Try to correlate with human judgment as closely as possible. You
will be given a context and a reply of the user U2 to the user U1 between the tags
<ORIGINAL_CONTEXT>. You will then be given a generated dialogue starting from the
short conversation, between the tags <GENERATED_DIALOGUE>. You will be given a set of
instructions between the tags <INSTRUCTIONS>. Your task is to evaluate the generated
dialogue (provided between the tags <GENERATED_DIALOGUE>) according to a specific
dimension (provided between the tags <DIMENSION_DEFINITION>), taking into account
the original context (provided between the tags <ORIGINAL_CONTEXT>). To evaluate the
dimension provided between the tags <DIMENSION_DEFINITION>, answer the questions provided
in <QUESTIONS>, strictly with 'yes' or 'no', following the instructions given in the tag
<INSTRUCTIONS> and the response format specified between the tags <ANSWER_FORMAT>.

<DIMENSION_DEFINITION>
{dimension_name} - {dimension_definition} </DIMENSION_DEFINITION>

<INSTRUCTIONS>
1. Carefully read the original context (between the tags <ORIGINAL_CONTEXT>) and the
generated dialogue (between the tags <GENERATED_DIALOGUE>).
2. Review the given questions associated with the dimension "{dimension_name}".
3. For each question between the tags <QUESTIONS>, answer strictly with 'yes' or 'no'.
4. Base your answers on clear reasoning, but **do not include the reasoning in the
output**.
5. Evaluate the generated dialogue, taking into account the short conversation.
6. Follow the answer format exactly as specified between the tags <ANSWER_FORMAT>.
</INSTRUCTIONS>

<ORIGINAL_CONTEXT>{original_turn} </ORIGINAL_CONTEXT>

<GENERATED_DIALOGUE>{all_turns} </GENERATED_DIALOGUE>

<QUESTIONS>{questions}</QUESTIONS>

<ANSWER_FORMAT>{answers_format} </ANSWER_FORMAT>

Provide your answers to each question following the specified Answer Format.
</TASK_OVERVIEW>
```

MP LLM-as-a-Judge Prompt

```
<TASK_OVERVIEW>
You are an automated judge model. You will receive:
- a short seed conversation between participants (tag: <SHORT_CONVERSATION>);
- a generated previous conversation (tag: <GENERATED_DIALOGUE>);
- a dimension definition (tag: <DIMENSION_DEFINITION>);
- a set of yes/no questions (tag: <QUESTIONS>);
- and an exact answer format (tag: <ANSWER_FORMAT>).

Your job: evaluate the GENERATED_DIALOGUE with respect to the DIMENSION_DEFINITION,
taking into account the SHORT_CONVERSATION. For each question in <QUESTIONS> answer
strictly with 'yes' or 'no' (lowercase). Do not output any additional text, explanation,
or chain-of-thought.

Answering rules:
- Read all provided content carefully.
- For each question, return exactly 'yes' or 'no' (lowercase) in the positions required by
<ANSWER_FORMAT>.
- Do NOT include reasoning, internal deliberation, or any extra characters (no
punctuation, no commentary).

Output:
- Produce only the content specified by <ANSWER_FORMAT> populated with 'yes'/'no' answers.

<DIMENSION_DEFINITION>
{dimension_name} - {dimension_definition}
</DIMENSION_DEFINITION>

<SHORT_CONVERSATION>{original_turn} </SHORT_CONVERSATION>

<GENERATED_DIALOGUE>{all_turns} </GENERATED_DIALOGUE>

<QUESTIONS>{questions}</QUESTIONS>

<ANSWER_FORMAT>{answer_format}
</ANSWER_FORMAT>
</TASK_OVERVIEW>
```

C. *Ironic-Response Affordance Detection Prompts*

CSC Response-Observed Reference Prompt ($C_{descriptive} + M_2^{U_2}$)

```
<TASK_OVERVIEW>
You have a pair of users, U1 and U2, involved in a dialogue.
The context, shown between the tags <CONTEXT>...</CONTEXT>, consists of a description of
the situation and the reply of U2.
Answer the following question, taking into account how U2 actually replied: Given this
dialogue history, would an ironic response be a plausible and natural continuation for
the upcoming respondent
The output must be 0 (not afforded) or 1 (afforded).
Provide your output answer between the tags <ANSWER></ANSWER>.
Provide your reasoning between the tags <REASONING></REASONING>.

<CONTEXT>{input}</CONTEXT>
</TASK_OVERVIEW>
```

CSC Single-Message Affordance Prompt ($M_2^{U_1}$)

```
<TASK_OVERVIEW>
You have a pair of users, U1 and U2, involved in a dialogue.
The context, shown between the tags <CONTEXT>...</CONTEXT>, consists of a message from U1,
to which U2 has to respond.
Answer the following question: Given this dialogue history, would an ironic response be a
plausible and natural continuation for the upcoming respondent?
```

```
The output must be 0 (not afforded) or 1 (afforded).
Provide your output answer between the tags <ANSWER></ANSWER>.
Provide your reasoning between the tags <REASONING></REASONING>.

<CONTEXT>{input}</CONTEXT>
</TASK_OVERVIEW>
```

CSC Multi-Turn Affordance Prompt ($M_0^{U_1} + M_0^{U_2} + \ldots + M_2^{U_1}$)

```
<TASK_OVERVIEW>
You have a pair of users, U1 and U2, involved in a dialogue.
The context, shown between the tags <CONTEXT>...</CONTEXT>, consists of two turns and a
last message from U1, to which U2 has to respond.
Answer the following question: Given this dialogue history, would an ironic response be a
plausible and natural continuation for the upcoming respondent?
The output must be 0 (not afforded) or 1 (afforded).
Provide your output answer between the tags <ANSWER></ANSWER>.
Provide your reasoning between the tags <REASONING></REASONING>.

<CONTEXT>{input}</CONTEXT>
</TASK_OVERVIEW>
```

MP Response-Observed Reference Prompt ($M_2^{U_1} + M_2^{U_2}$)

```
<TASK_OVERVIEW>
You have a pair of users, U1 and U2, involved in a dialogue.
The context, shown between the tags <CONTEXT>...</CONTEXT>, consists of a pair of
messages exchanged on social media.
Answer the following question, taking into account how U2 actually replied: Given this
dialogue history, would an ironic response be a plausible and natural continuation for
the upcoming respondent?
The output must be 0 (not afforded) or 1 (afforded).
Provide your output answer between the tags <ANSWER></ANSWER>.
Provide your reasoning between the tags <REASONING></REASONING>.

<CONTEXT>{input}</CONTEXT>
</TASK_OVERVIEW>
```

MP Single-Message Affordance Prompt ($M_2^{U_1}$)

```
<TASK_OVERVIEW>
You have a pair of users, U1 and U2, involved in a dialogue.
The context, shown between the tags <CONTEXT>...</CONTEXT>, consists of a message from U1
on social media, to which U2 has to respond.
Answer the following question: Given this dialogue history, would an ironic response be a
plausible and natural continuation for the upcoming respondent?
The output must be 0 (not afforded) or 1 (afforded).
Provide your output answer between the tags <ANSWER></ANSWER>.
Provide your reasoning between the tags <REASONING></REASONING>.

<CONTEXT>{input}</CONTEXT>
</TASK_OVERVIEW>
```

MP Multi-Turn Affordance Prompt ($M_0^{U_1} + M_0^{U_2} + \ldots + M_2^{U_1}$)

```
<TASK_OVERVIEW>
You have a pair of users, U1 and U2, involved in a dialogue.
The context, shown between the tags <CONTEXT>...</CONTEXT>, consists of two turns and a
last message from U1, to which U2 has to respond.
Answer the following question: Given this dialogue history, would an ironic response be a
plausible and natural continuation for the upcoming respondent?
The output must be 0 (not afforded) or 1 (afforded).
```

```
Provide your output answer between the tags <ANSWER></ANSWER>.
Provide your reasoning between the tags <REASONING></REASONING>.

<CONTEXT>{input}</CONTEXT>
</TASK_OVERVIEW>
```

D. *Perspectivist Affordance Detection Prompts*

Trait-Based Affordance Prompt

```
<TASK_OVERVIEW>
You are {trait}. Use this perspective when making the following decision.
You have a pair of users, U1 and U2, involved in a dialogue. The context contains two
turns and a last message from U1, to which U2 has to respond.
Given this dialogue history, would an ironic response be a plausible and natural
continuation for the upcoming respondent?
The output must be 0 (not afforded) or 1 (afforded).
Provide your output answer between the tags <ANSWER></ANSWER>.

<CONVERSATION>{input}</CONVERSATION>
</TASK_OVERVIEW>
```

Relation-Based Affordance Prompt

```
<TASK_OVERVIEW>
You have a pair of users, U1 and U2, involved in a dialogue. You know that U1 and U2 are
{relation}.
The context contains two turns and a last message from U1, to which U2 has to respond.
Given this dialogue history, would an ironic response be a plausible and natural
continuation for the upcoming respondent?
The output must be 0 (not afforded) or 1 (afforded).
Provide your output answer between the tags <ANSWER></ANSWER>.

<CONVERSATION>{input}</CONVERSATION>
</TASK_OVERVIEW>
```